



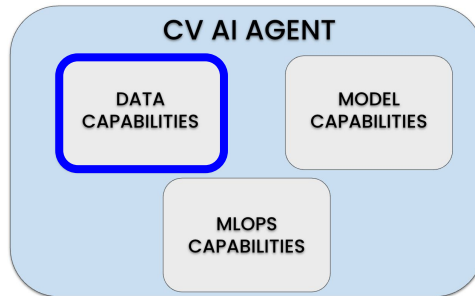
AgenticVision

User Journey of the agentic CV platform

HTX User Requirements/Capabilities (Data-specific only)

This user journey maps out how users interact with the agentic data management system. The system is capable of performing the following :

1. **Data preprocessing** : Reduces the current 70% of time HTX currently spends on data processing, curation, and cleaning
2. **Data Annotation** : Resolves inconsistent data standardization and labeling issues (e.g., swimmer definitions, boat classifications)
3. **Data Quality** : Automates quality assessment of visual assets
4. **Data Augmentation and generation** : Augments vision datasets for data balancing and model generalization capabilities.
5. **Search** : Enables efficient search and retrieval across video repositories



User journey

Broken down into three main aspects :

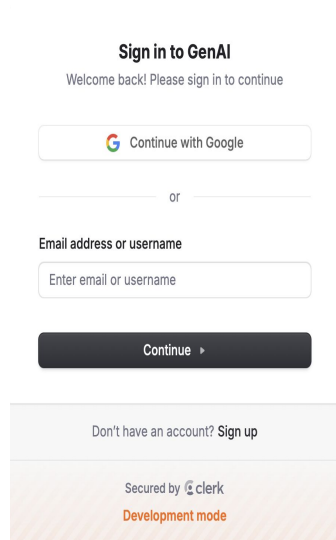
1. Initial setup and configuration
2. Automatic mode detection and tasks generation
3. Execution flows

Phase 1 : Initial Setup & Configuration


1. System onboarding - login

User Actions: logging in

- Access a web interface interface hosted on a server, show login page and authenticate using clerk (left image)
- Create profile and complete setup if account not created using clerk (right image)



Sign in to GenAI
Welcome back! Please sign in to continue


 Continue with Google

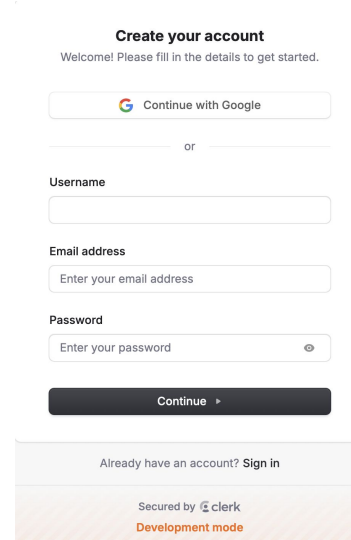
or

Email address or username


Continue »

Don't have an account? [Sign up](#)

Secured by  clerk
Development mode



Create your account
Welcome! Please fill in the details to get started.

 Continue with Google

or

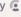
Username

Email address

Password

Continue »

Already have an account? [Sign in](#)

Secured by  clerk
Development mode

1. System onboarding - Projects

User actions: Creating new project

- Show a list of all projects
- A project contains dataset and the associated models
- Datasets and models will be versioned within each project
- User Clicks on “New Project” to create a project

Projects

The screenshot displays the RoboFlow Projects interface. On the left, a list of projects is shown, each with a thumbnail, a title, a description, and a status. A purple button labeled '+ New Project' is located at the top right of the list. An arrow points from this button to a modal window on the right titled 'Let's create your project.' The modal contains a 'Project Name' field with a placeholder text 'E.g., 'Dog Breeds' or 'Car Models' or 'Text Finder'.', a 'Tags' dropdown menu, a 'Description' field with a placeholder text 'E.g., 'dogs' or 'cars' or 'words'.', and a 'New Public Project' link.

Projects

+ New Project

Object Detection
🌐 **Yard**
Edited 20 days ago
Public • 220 Images

Object Detection
🌐 **Yard Management Syst...**
Edited 5 months ago
Public • 186 Images • 3 Models

Object Detection
🌐 **Aerial Maritime OBB**
Edited 6 months ago
Public • 102 Images

Object Detection
🌐 **Mask Wearing**
Edited 8 months ago
Public • 3530 Images • 5 Models

Let's create your project.
RoboFlow Universe Projects > [New Public Project](#)

Project Name
E.g., 'Dog Breeds' or 'Car Models' or 'Text Finder'.

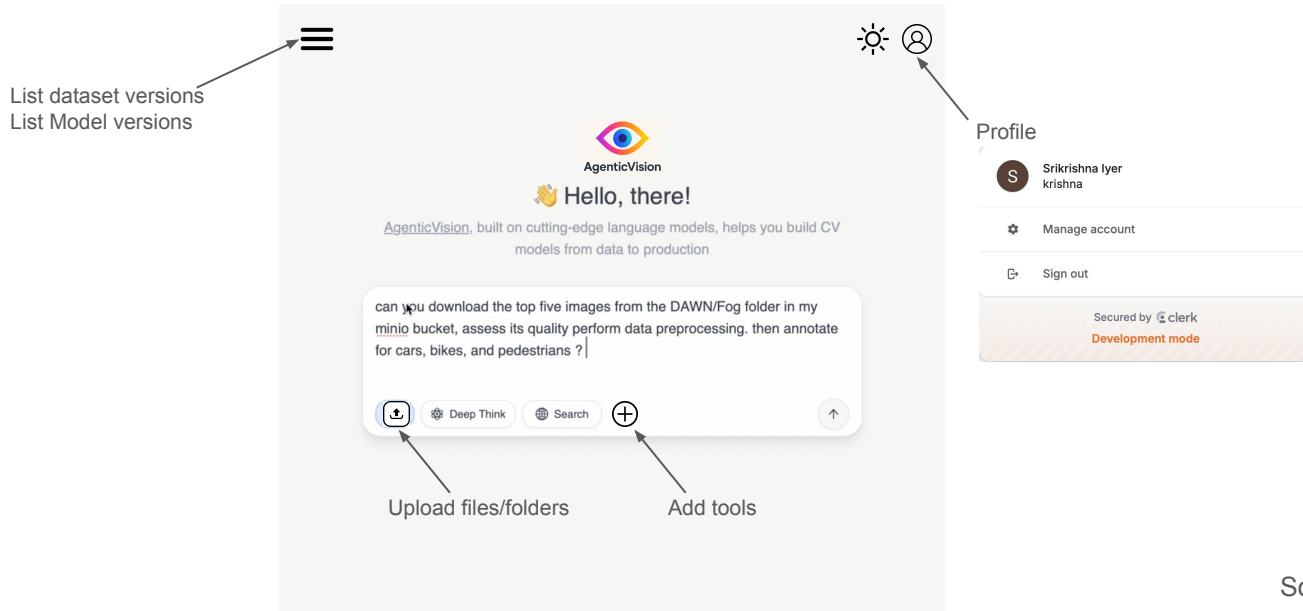
Tags
▼

Description
E.g., 'dogs' or 'cars' or 'words'.

1. System onboarding - Chat window

System Response: First page

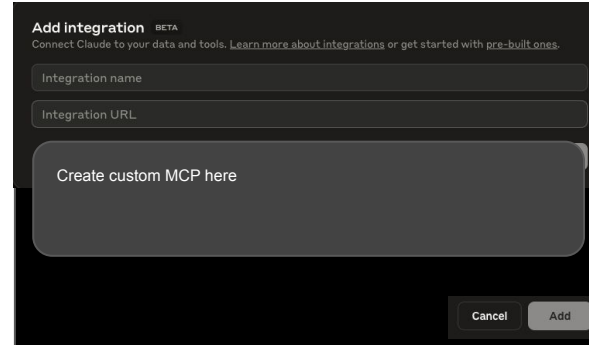
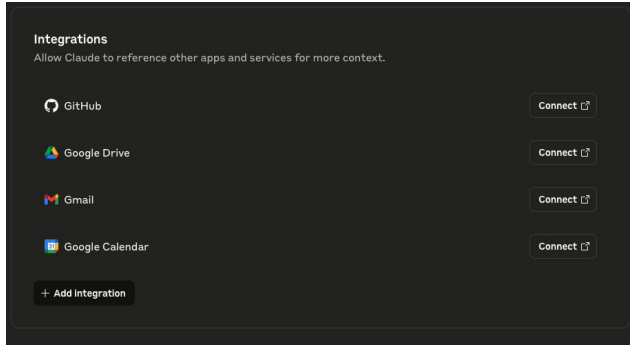
- Show welcome interface with a textbox.
- Buttons from left to right : upload folder, Deep think (TBD), Search (TBD), tools (show existing tools, modify tools)
- Clicking the “+” icon, takes you to a list of available tools and how you can add/modify/delete them



1. System onboarding - Adding custom tools

User Actions: Adding tools

- Clicking on the “+” icon
- Add and modify MCP tools



1. System onboarding - Upload files/folders

User Actions: Uploading data

- Clicking on the “upload” icon to open a pop up menu
- If uploading videos, the system should automatically detect .mp4 files and pop up an additional menu to select fps (right image)

The 'Upload' dialog box features a title bar with an upward arrow icon and the word 'Upload'. Below the title bar, there is a 'Batch Name:' label followed by a text input field containing the word 'Data'. To the right of this is a 'Tags:' label with a circular icon containing a plus sign, followed by a text input field with the placeholder text 'Search or add tags for images...'. In the center of the dialog is a large dashed rectangular area. Above this area is a purple circle containing a white upward arrow. Below the arrow is the text 'Drag and drop file(s) to upload, or:'. Underneath this text are two buttons: 'Select File(s)' and 'Select Folder'. Below these buttons is a section titled 'Supported Formats' which contains four categories: 'Images' (with a camera icon), 'Annotations' (with a square icon), 'Videos' (with a film strip icon), and 'PDFs' (with a document icon). Each category lists supported file extensions: Images (.jpg, .png, .bmp, .webp, .avif), Annotations (in 26 formats), Videos (.mov, .mp4), and PDFs (.pdf). A note at the bottom of the supported formats section states 'Max size of 20MB and 16,384 pixels per dimension.' At the bottom right of the dialog is a purple button labeled 'Skip'.

If .mp4
detected

This dialog box is titled 'How often should we sample this video?'. It displays a video player interface showing a screen recording. Below the video player, the text 'Screen Recording 2025-05-13 at 10.15.03.mov (17.1s)' is shown. Below the video player is a horizontal timeline with a yellow highlight. Below the timeline are three radio button options: '60 frames/second', '1 frame/second' (which is selected), and '1 frame every 60 seconds'. Below these options is a horizontal slider with a blue circle in the middle. Below the slider is the text 'Output Size: 18 images'. At the bottom of the dialog are three buttons: 'Apply to All', 'Apply', and 'Choose Frame Rate'.

Phase 2 : automatic Mode detection and task generation

2. Intelligent mode detection

System Action: Mode detection

A planner agent (reasoning LLM eg. o3) will determine which type of mode is triggered based on user input :

a. Fully Autonomous Triggers:

- Open-ended natural language requests
- Minimal technical specifications
- General goals without specific methodologies
- Examples: "Collect Singapore traffic datasets", "Label these images", "Clean up this dataset", "Make these photos better", "Annotate the images"

b. User Specified Triggers:

- Detailed technical requirements
- Specific tool mentions or parameters
- Explicit pipeline steps mentioned
- Examples: "Use YOLO for detection with confidence threshold 0.85", "Apply Gaussian blur with sigma=2.1"

c. Hybrid Triggers:

- Partial specifications with room for optimization
- Requests for collaboration or recommendations
- Mixed technical and natural language elements
- Quality requirements with flexible implementation
- Examples: "Improve quality while maintaining file sizes under 2MB", "Denoise the images but also check if they can be improved further", "apply cutmix, resize images to 128x128 but also try other data augmentation techniques"

2. Task generation

Phase 3 : Execution Flows

Backend : Workflow Orchestration

The system should support a “**combination**” of two workflows :

1. **Static Workflows (RPA):** Predefined sequences designed for repetitive tasks.

Tool	Data Ingestion	Data Processing	Data Annotation	Augmentation	Data Quality Assessment	Automated CI/CD	Method
Roboflow	✓ (image uploads APIs/CLI)	✓ (auto preproc pipeline)	✓ (built-in labeling)	✓ (resize, flip, color augment)	🔧 (via scripts + validation API)	🔧 (via API + CI pipelines)	low-code + APIs/CLI
Airflow	✓ (operators for storage, HTTP, etc.)	✓ (Python tasks/operators)	✗	🔧 (scripted within DAGs)	🔧 (via hooks, custom tasks)	✓ (native DAG scheduling + triggers)	code (Python)
Kubeflow	✓ (via Pipelines/KFServing)	✓ (components, pipelines)	✗ (needs script or plugin)	🔧 (script in components)	🔧 (via MLMD, metrics)	✓ (Pipeline orchestration + CI/CD hooks using kubernetes)	code (Python, YAML)

2. **Agentic Workflows:** Dynamic workflows capable of adapting to different conditions. For some use cases, calling static workflows as an API or python function via MCP.


Tool	Long-running Tasks	System Integration	Human-in-the-Loop	Method
LangGraph	✅ Yes	✅ Yes (LangServe)	✅ Yes (supports human-in-loop)	Code (Python, graph-based orchestrator)
Dify	✅ Yes (workflow engine, celery powered)	✅ Yes (connects LLMs & APIs via UI & API)	❌ No	Low-code (UI workflows + prompt IDE)
n8n	✅ Yes (via external async, supports workflows)	✅ Yes (800+ connectors, HTTP node)	✅ via Human-in-the-Loop node (paid plan)	Low-code / No-code + JS/TS scripting
CrewAI	✅ Yes (Flow engine/Celery-powered)	❌ No (no built-in system integrations; but tool/plugin support)	❌ Partial (via plugin/UI, but limited interactivity)	Code (Python) + some no-code Studio

Backend : Workflow Orchestration

Components

Each workflow may include the following components:

- **Input:** Data sources or user interactions that initiate the workflow.
- **LLM/VLM:** Language/Vision-language models for understanding and processing tasks.
- **Memory:** Context retention to support long-term reasoning.
- **MCP Tool integration:** External APIs, functions or static workflows that perform specific actions.
- **Operator:** Logic that controls the workflow's flow and decision-making.
- **Output:** Final results or actions produced by the workflow.




3.1 : Execution Flows ~ Data collection

System Actions : Data collection

- Mode dependent data collection
 - Fully autonomous
 - Open-ended request to collect datasets
 - Use browser-use agent that uses a headless browser to emulate mouse clicks to download datasets
 - The system combines datasets and their annotations from multiple sources. Eg. Huggingface, Kaggle
 - User specified
 - Upload data manually within the UI
 - Add MCP tool connection to external storage. Eg. MinIO, S3
 - System makes sure annotation formats are aligned to user request.
 - Hybrid
 - Upload data manually within the UI
 - User can also request additional data request and combine with the uploaded data. For eg. “Collect 10K samples of hazy traffic data to augment my uploaded dataset.”



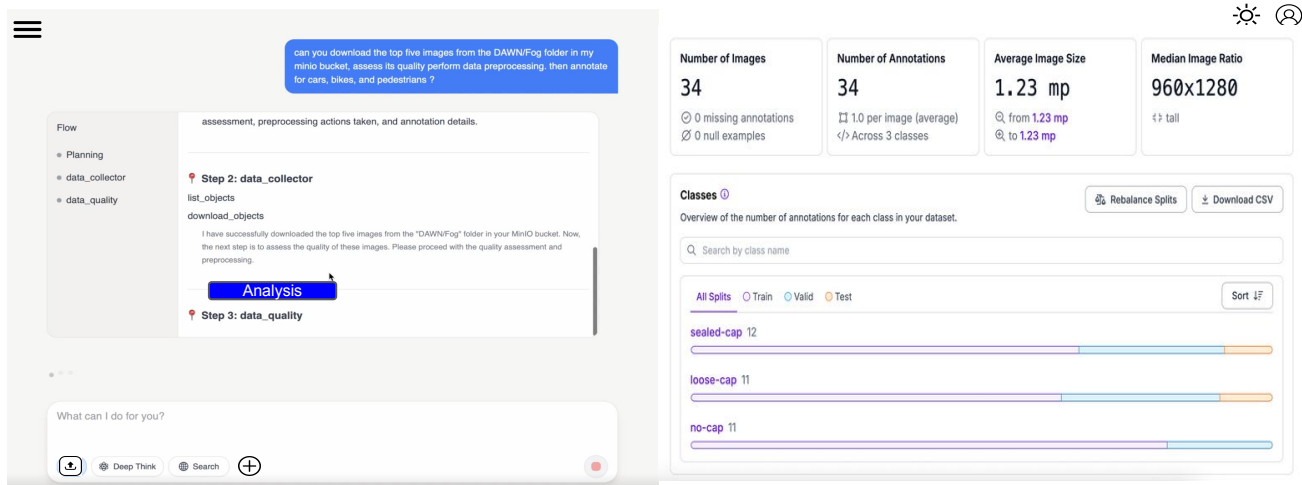
Example browser agent for automatic data collection



3.1 : Execution Flows ~ Data Collection

System Response: Data analysis

- Split screen view :
 - The left panel shows the agent's execution flow, along with the agent's intermediate thoughts and “artefacts” (optional). The artefacts will maximize the right panel.
 - Right panel shows
 - Output artefacts : Only for visualizing the intermediate outputs of each step. Show output previews of 5 randomly selected images/videos
 - Interactable human in the loop components : To approve/deny/input the workflow intermediate outputs
- After data is collected, we could show metadata of the datasets depending on the availability of annotations.
- If annotations exist, we show class splits (as shown below)
- Further analysis could include “image dimensions”, “annotation heat map”, “Histogram count of objects by image”. Refer [here](#) for examples.



3.2 : Execution Flows ~ Data quality

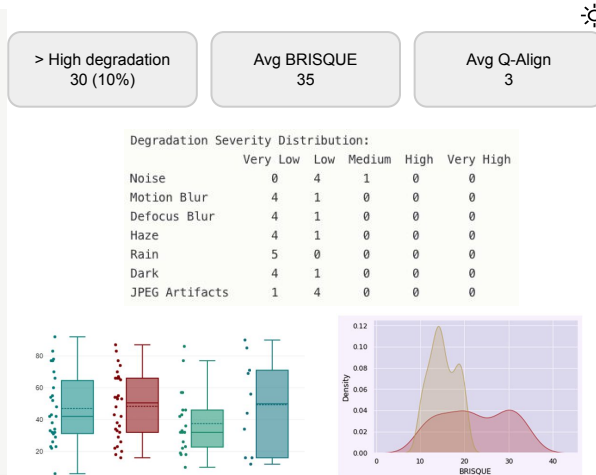
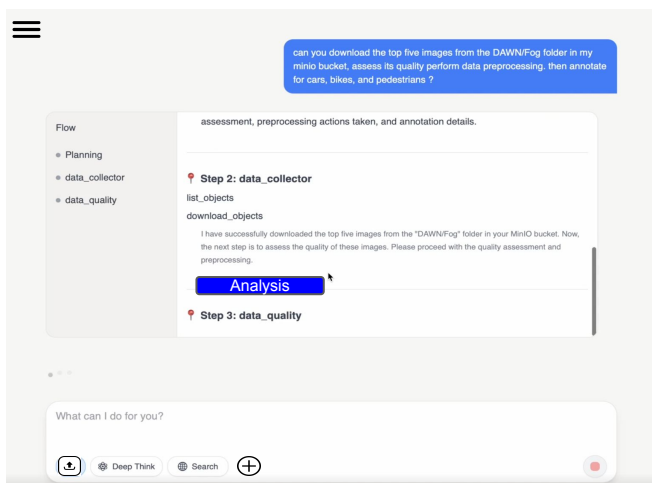
System Actions

- Fully autonomous
 - Open-ended request to perform quality analysis. Eg. “Analyse and detect degradations in the dataset”
 - No-reference IQA + VLM driven degradation detection. Refer [here](#) for more info.
- User specified
 - User specifies what type and metric to assess quality of images. Eg. “Detect hazy images and return a quality score using BRISQUE”
- Hybrid
 - User specifies what type and metric to assess quality of images.
 - But also asks agent to detect other degradations if any. Eg. “Detect hazy images and return a quality score using BRISQUE. Also detect if any other degradations exist.”

3.2 : Execution Flows ~ Data quality

System Response

- Display analysis on the data quality assessment :
 - Severity distribution : Count of severity by degradation type
 - Metric distribution : Box plots, histograms, violin plots, scatter plots by metric
- Further analysis could include “image dimensions”, “annotation heat map”, “Histogram count of objects by image”. Refer [here](#) for examples.



3.2 : Execution Flows ~ Data preprocessor

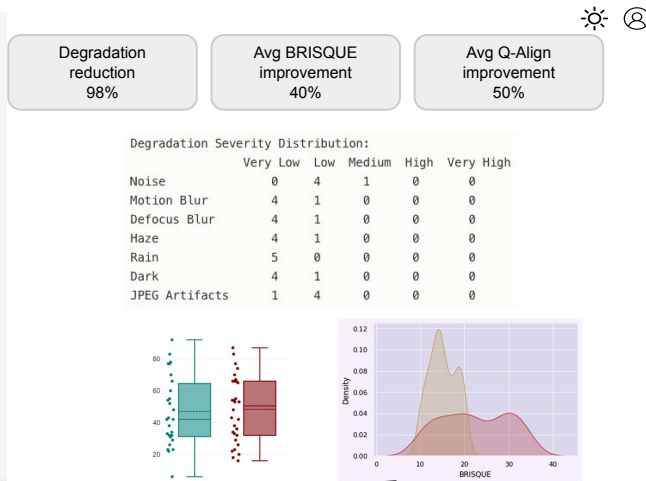
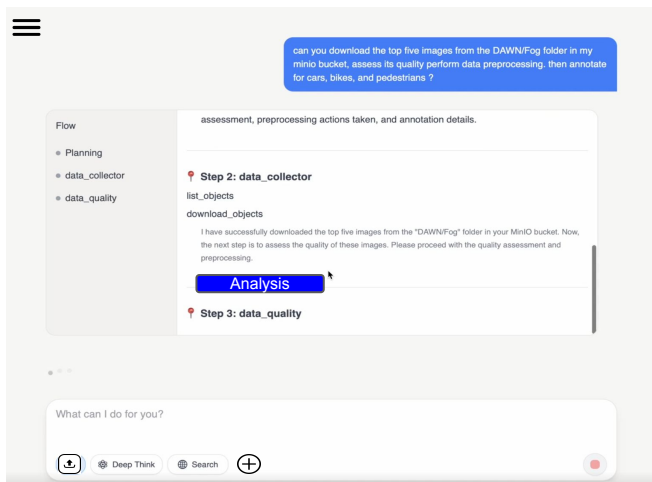
System Actions : Data preprocessor

- Fully autonomous
 - Open-ended request to perform preprocessing. Eg. “Preprocess the images”
 - The system should first perform autonomous data quality check.
 - Then divide the images into batches (by degradation type)
 - Then determine the preprocessing steps for each batch. Refer [here](#) for workflow.
- User specified
 - User specifies what type and metric to assess quality of images. Eg. “Detect hazy images and return a quality score using BRISQUE”
- Hybrid
 - User specifies what type and metric to assess quality of images.
 - But also asks agent to detect other degradations if any. Eg. “Detect hazy images and return a quality score using BRISQUE. Also detect if any other degradations exist.”

3.2 : Execution Flows ~ Data preprocessor

System Response: Data preprocessor

- Display analysis on before and after data processing
 - Stats showing % drop in degradation, BRISQUE scores, Q-Align scores
 - Severity distribution : Count of severity by degradation type
 - Metric distribution side by side comparison : Box plots, histograms, violin plots, scatter plots by metric



Side by side comparison of quality metrics before and after processing

3.2 : Execution Flows ~ Data Annotation

System Actions

- Fully autonomous
 -
- User specified
 -
- Hybrid
 -

3.2 : Execution Flows ~ Data Annotation

System Response

3.2 : Execution Flows ~ Data Augmentation

System Actions

- Fully autonomous
 -
- User specified
 -
- Hybrid
 -

3.2 : Execution Flows ~ Data Augmentation

System Response

3.2 : Execution Flows ~ Data Search

System Actions

- Fully autonomous
 -
- User specified
 -
- Hybrid
 -

3.2 : Execution Flows ~ Data Search

System Response

Appendix

Agentic AutoML : A review

Table 1. Comparison between *AutoML-Agent* and existing LLM-based frameworks.

Framework	Key Functionality					
	Planning	Verification	Full Pipeline	Task-Agnostic	Training-Free Search	With Retrieval
AutoML-GPT (Zhang et al., 2023)	×	×	×	✓	✓	×
Prompt2Model (Viswanathan et al., 2023)	×	×	✓	×	×	✓
HuggingGPT (Shen et al., 2023)	✓	×	×	✓	✓	✓
CAAFE (Hollmann et al., 2023b)	×	✓	×	×	×	×
MLCopilot (Zhang et al., 2024a)	×	×	×	✓	✓	×
AgentHPO (Liu et al., 2025)	✓	✓	×	✓	×	×
Data Interpreter (Hong et al., 2024a)	✓	✓	×	✓	×	×
DS-Agent (Guo et al., 2024a)	✓	✓	×	✓	×	✓
SELA (Chi et al., 2024)	✓	✓	×	✓	×	×
Agent K (Grosnit et al., 2024)	✓	✓	×	✓	×	✓
AutoMMLab (Yang et al., 2025)	×	✓	✓	×	×	×
<i>AutoML-Agent</i>	✓	✓	✓	✓	✓	✓

AutoMMLab: Automatically Generating Deployable Models from Language Instructions for Computer Vision Tasks

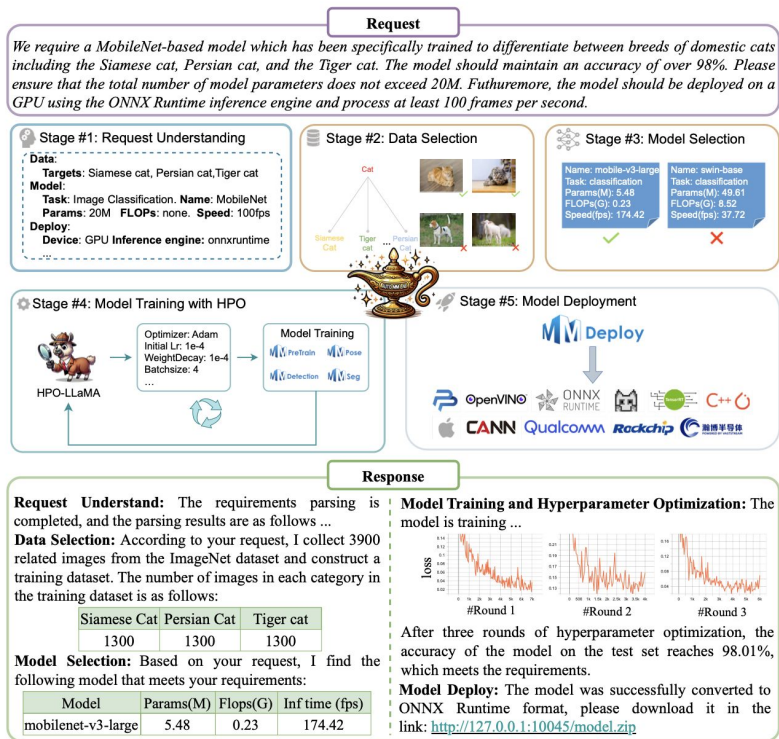


Figure 2: Overview of AutoMMLab. The workflow of AutoMMLab consists of five stages. **Request understanding:** Parse the language requests into formatted configuration. **Data selection:** Select appropriate training data from the dataset zoo. **Model selection:** Select the optimal model from the model zoo. **Model training with HPO:** Train the model and optimize the hyperparameters. **Model deployment:** Convert the model into a package compatible with the deployment environments.

The authors release a benchmark called LAMP for autoML tasks to evaluate :

1. Request Understanding :
 - a. Key-level accuracy calculates the average accuracy of each key-value pair.
 - b. Req-Level accuracy calculates the accuracy of understanding the entire request.
2. Hyperparameter Optimization : Propose LLM-based optimization (seems expensive to me, better to use Bayesian or other simpler methods that achieve similar performance)
3. End-to-end Evaluation : Pass/Fail

Hyperparameter optimization using LLaMA

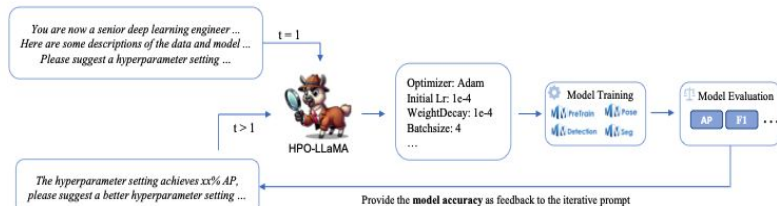


Figure 4: Overview of HPO-LLaMA. At the initial step ($t = 1$), HPO-LLaMA proposes a hyperparameter configuration based on the description of model and task. Model training is then executed and the training results are passed back to HPO-LLaMA via a text prompt for further rounds ($t > 1$).

AutoMMLab : Results

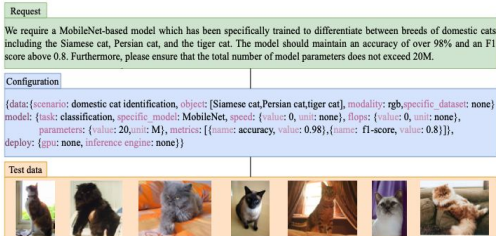


Figure 3: Example of LAMP dataset.

- RU-LLaMA finetuned on (request, configuration) pairs generated by GPT-4
- HPO-LLaMA finetuned on randomly selected 8000 (request, hyperparameters, performance) triplets.

Model	Key-Level			Req-Level
	Item	List	Total	
LLaMA2-7B-Chat	85.71	50.00	77.78	0
PaLM2	96.79	88.13	94.86	63.75
GPT-3.5-turbo	96.43	95.63	96.25	72.50
GPT-4	97.50	93.13	96.53	80.00
RU-LLaMA	98.57	96.88	98.20	86.25

Table 1: Evaluation of request understanding (RU). Best results are marked in bold.




Figure 5: HPO results of HPO-LLaMA and random sampling baselines on four tasks : (a) image classification, (b) object detection, (c) semantic segmentation and (d) keypoint detection. HPO-LLaMA demonstrates significantly higher efficiency.

Model	#R	Cls.	Det.	Seg.	Kpt.
BayesianRF	5	0.618±0.287	0.291±0.211	0.847±0.044	0.069±0.136
BayesianGP	5	0.761±0.264	0.280±0.208	0.848±0.041	0.081±0.150
LLaMA2-7B	1	0.839±0.213	0.128±0.164	0.291±0.409	0±0
PaLM2	1	0.964±0.056	0.367±0.196	0.845±0.067	0.719±0.079
GPT-3.5-turbo	1	0.849±0.214	0.364±0.194	0.852±0.044	0.204±0.160
GPT-4	1	0.861±0.188	0.434±0.147	0.803±0.194	0.096±0.158
HPO-LLaMA	1	0.975±0.028	0.435±0.148	0.854±0.042	0.728±0.051
HPO-LLaMA	3	0.983±0.020	0.440±0.150	0.856±0.043	0.738±0.053

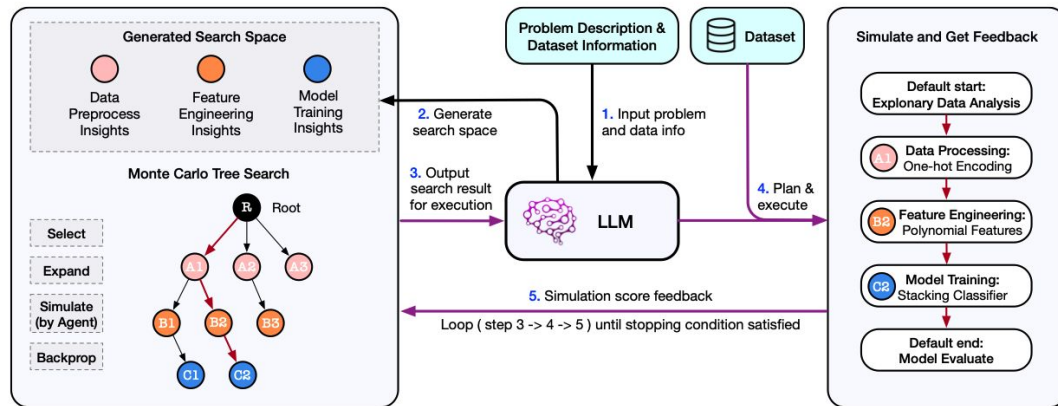
Table 2: Evaluation of HPO. #R means the number of iteration rounds. The mean and standard deviation on for task are exhibited. Best results are marked in bold, second best results are underlined.

The models achieve similar results using random search, only slower (by 6-7 iterations). LLM optimization is an overkill and quite expensive if it doesn't converge.

RU	HPO	Cls.	Det.	Seg.	Kpt.	Total
LLaMA2-7B	LLaMA2-7B	0	0	0	0	0
PaLM2	PaLM2	14	25	27	15	81
GPT-3.5-turbo	GPT-3.5-turbo	24	24	25	11	84
GPT-4	GPT-4	17	27	29	14	87
RU-LLaMA	HPO-LLaMA	31	31	32	18	112

Table 3: End-to-end evaluation on LAMP benchmark. The assessment is based on a grading system scoring from '0' to '2', where '0' denotes 'total failure', '1' denotes 'workable model', 2 denotes 'perfect model'. The full score for each task is 40, and the total full score is 160.

SELA: TREE-SEARCH ENHANCED LLM AGENTS FOR AUTOMATED MACHINE LEARNING



RESULTS

$$NS(s_{raw}) = \begin{cases} \frac{1}{1 + \log(1 + s_{raw})} & \text{if the metric is RMSE.} \\ s_{raw} & \text{otherwise.} \end{cases}$$

Figure 2: SELA's pipeline operates as follows: The system begins by inputting the problem description and dataset information into the LLM, which generates a search space of potential solutions, encompassing data preprocessing, feature engineering, and model training. The search module, powered by Monte Carlo Tree Search (MCTS), explores this space by selecting, expanding, and simulating potential configurations. The LLM agent then simulates the selected configuration by planning, coding, and executing the experiment. Feedback from the simulation is fed back into the search module, where it is used in the backpropagation step to refine future searches. This iterative process continues until a predefined stopping criterion is met, resulting in an optimized experimental pipeline.

DS-Agent: Automated Data Science by Empowering Large Language Models with Case-Based Reasoning

What is Case based reasoning ?

Retrieve similar past problems, reusing their solutions for the current problem, evaluating the effectiveness, revising the solution, and retaining successful solutions.

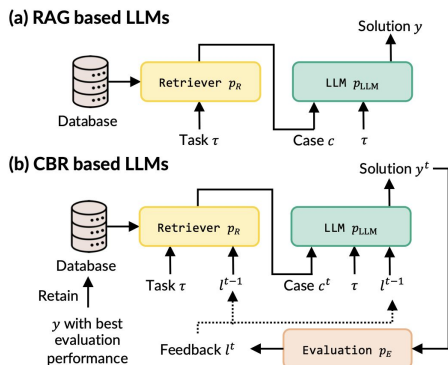


Figure 2. Comparison between (a) RAG based LLMs and (b) CBR based LLMs.

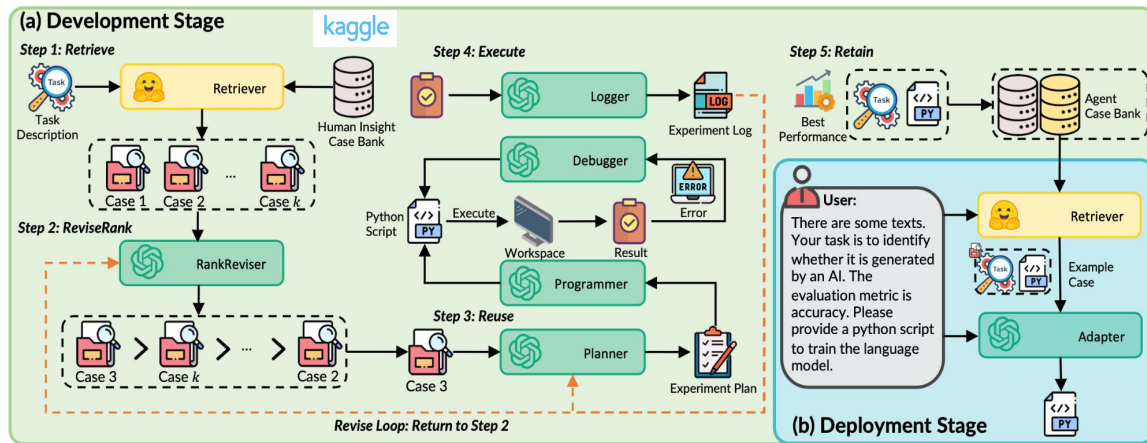


Figure 3. The diagram of DS-Agent. (a) **Development Stage**: DS-Agent structures an automatic iteration pipeline to build and revise the model based on execution feedback. (b) **Deployment Stage**: DS-Agent adapts past successful solutions for code generation.

DS-Agent: Results

Table 3. Mean rank w.r.t. task-specific evaluation metric results on 18 data science tasks in the deployment stage. Results are reported over 10 repetitive runs. Best performances are highlighted in bold, and second best performances are underlined.

		JS	HR	BPP	WR	DAG	BQ	TFC	WTH	ELE	SRC	UGL	HB	CA	CS	MH	SS	CO	SD	Avg
Mixtral-8x7b-Instruct	Zero-shot	37.0	35.0	35.0	31.0	35.0	32.0	29.0	32.0	30.0	44.0	54.0	46.0	73.1	66.6	65.8	63.6	33.7	72.0	45.3
	One-shot	35.2	35.0	32.2	31.0	35.0	29.1	29.0	32.0	30.0	36.5	47.1	46.0	50.1	53.1	51.2	51.1	23.6	61.5	39.4
	DS-Agent	37.0	35.0	35.0	31.0	35.0	32.0	29.0	32.0	30.0	<u>20.1</u>	16.4	38.5	<u>25.3</u>	54.5	53.7	53.9	32.2	47.6	35.5
GPT-3.5	Zero-shot	21.7	35.0	30.1	28.6	27.1	28.3	27.1	29.1	28.1	33.1	48.4	21.4	29.0	35.3	28.8	35.7	25.2	42.3	30.8
	One-shot	27.6	25.8	27.6	25.6	34.6	23.0	20.8	29.1	27.0	35.7	48.4	21.1	27.1	50.5	58.4	57.5	33.9	56.4	35.0
	DS-Agent	6.0	<u>22.6</u>	15.0	<u>20.6</u>	<u>15.1</u>	13.1	<u>17.3</u>	<u>13.4</u>	<u>14.4</u>	20.0	<u>13.0</u>	23.0	29.0	<u>19.3</u>	7.6	2.0	37.0	<u>19.5</u>	17.1
GPT-4	Zero-shot	36.7	31.8	35.0	29.0	29.4	32.0	29.0	32.0	30.0	37.3	45.7	33.6	1.0	15.3	23.2	17.9	28.3	20.1	28.2
	One-shot	35.1	24.4	13.8	26.6	29.6	28.8	23.1	30.1	26.6	26.7	41.6	36.7	29.7	21.9	35.3	28.9	<u>21.4</u>	23.2	28.0
	DS-Agent	<u>18.6</u>	1.0	<u>14.6</u>	5.2	6.2	<u>18.8</u>	15.7	6.3	8.1	20.0	11.4	<u>21.2</u>	1.0	32.6	<u>14.5</u>	<u>8.2</u>	13.0	12.4	12.7

Table 2. Ablation results in terms of average best rank over 12 development tasks. Results are reported over five repetitive trials.

GPT-4	Average Best Rank
DS-Agent	2.08
DS-Agent w/o ReviseRank	2.58
DS-Agent w/o CBR	3.41

Improves over past kaggle human experts




Figure 6. Further analyses on DS-Agent in the deployment stage. (a) Performance difference of DS-Agent learning from past successful experiences or textual human insights. (b) Hyper-parameter study on varying number of example case in DS-Agent with GPT-3.5.

AutoML-Agent: A Multi-Agent LLM Framework for Full-Pipeline AutoML

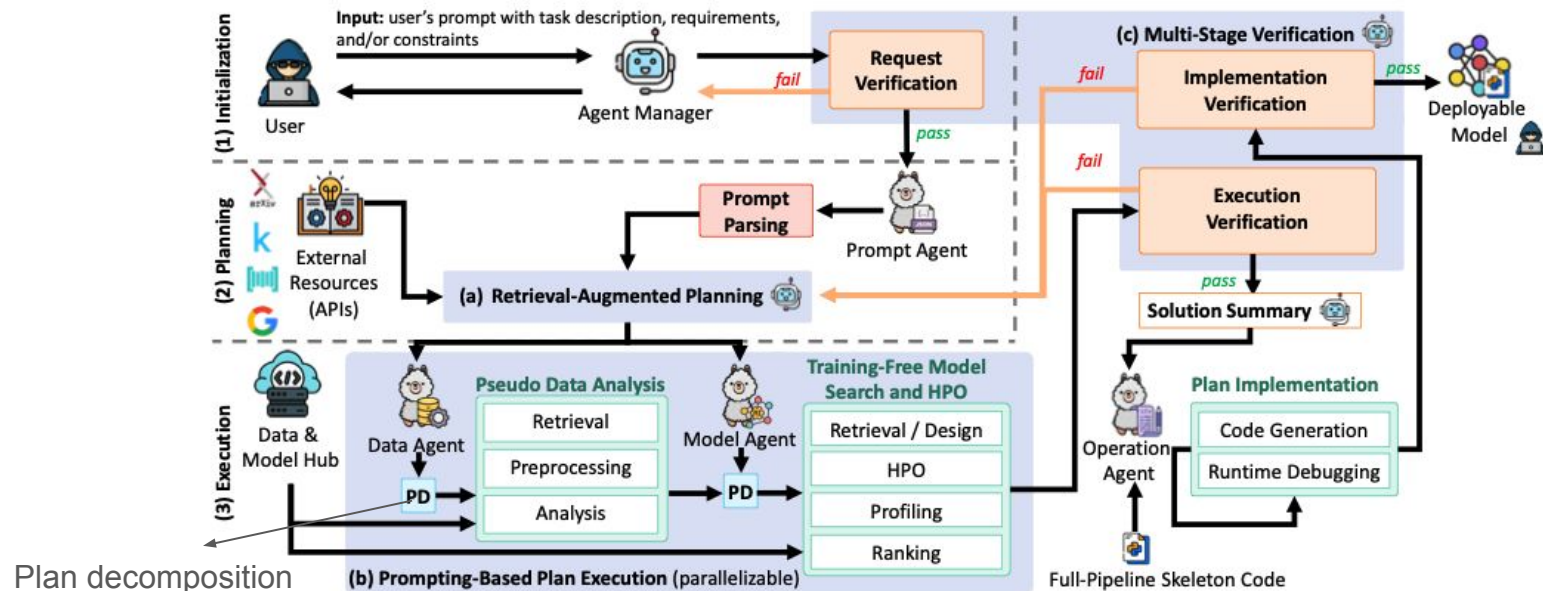


Figure 2. Overview of our AutoML-Agent framework. (1) **Initialization** stage aims to receive a valid user instruction using request verification. (2) **Planning** stage focuses on extracting ML related information by parsing the user instruction into a standardized form, and uses it to devise plans accordingly. (3) **Execution** stage executes each action given by the devised plans. Finally, based on the best execution results, AutoML-Agent outputs codes containing deployable model to the user.

AutoML-Agent : Results

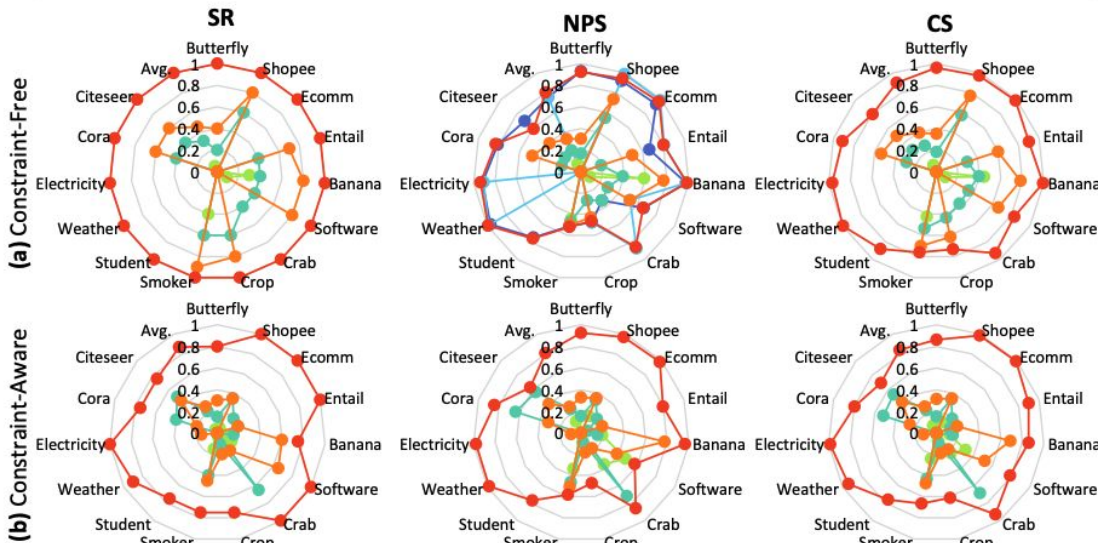


Figure 4. Performance comparison across all datasets using the SR, NPS, and CS metrics under (a) constraint-free and (b) constraint-aware settings. Higher scores indicate better results.



Figure 7. Average time and monetary cost breakdown.

Challenges of existing methods

Data challenges

- Lack of Automated data sourcing
- Lack of standardization of varied datasets (and annotation formats)
- Complex preprocessing and feature engineering for CV tasks
- Do not cover all CV tasks
- Skipped annotation, data augmentation entirely
- No data quality based verification step
- Preprocessing is fixed, does not depend on data distribution, quality and task

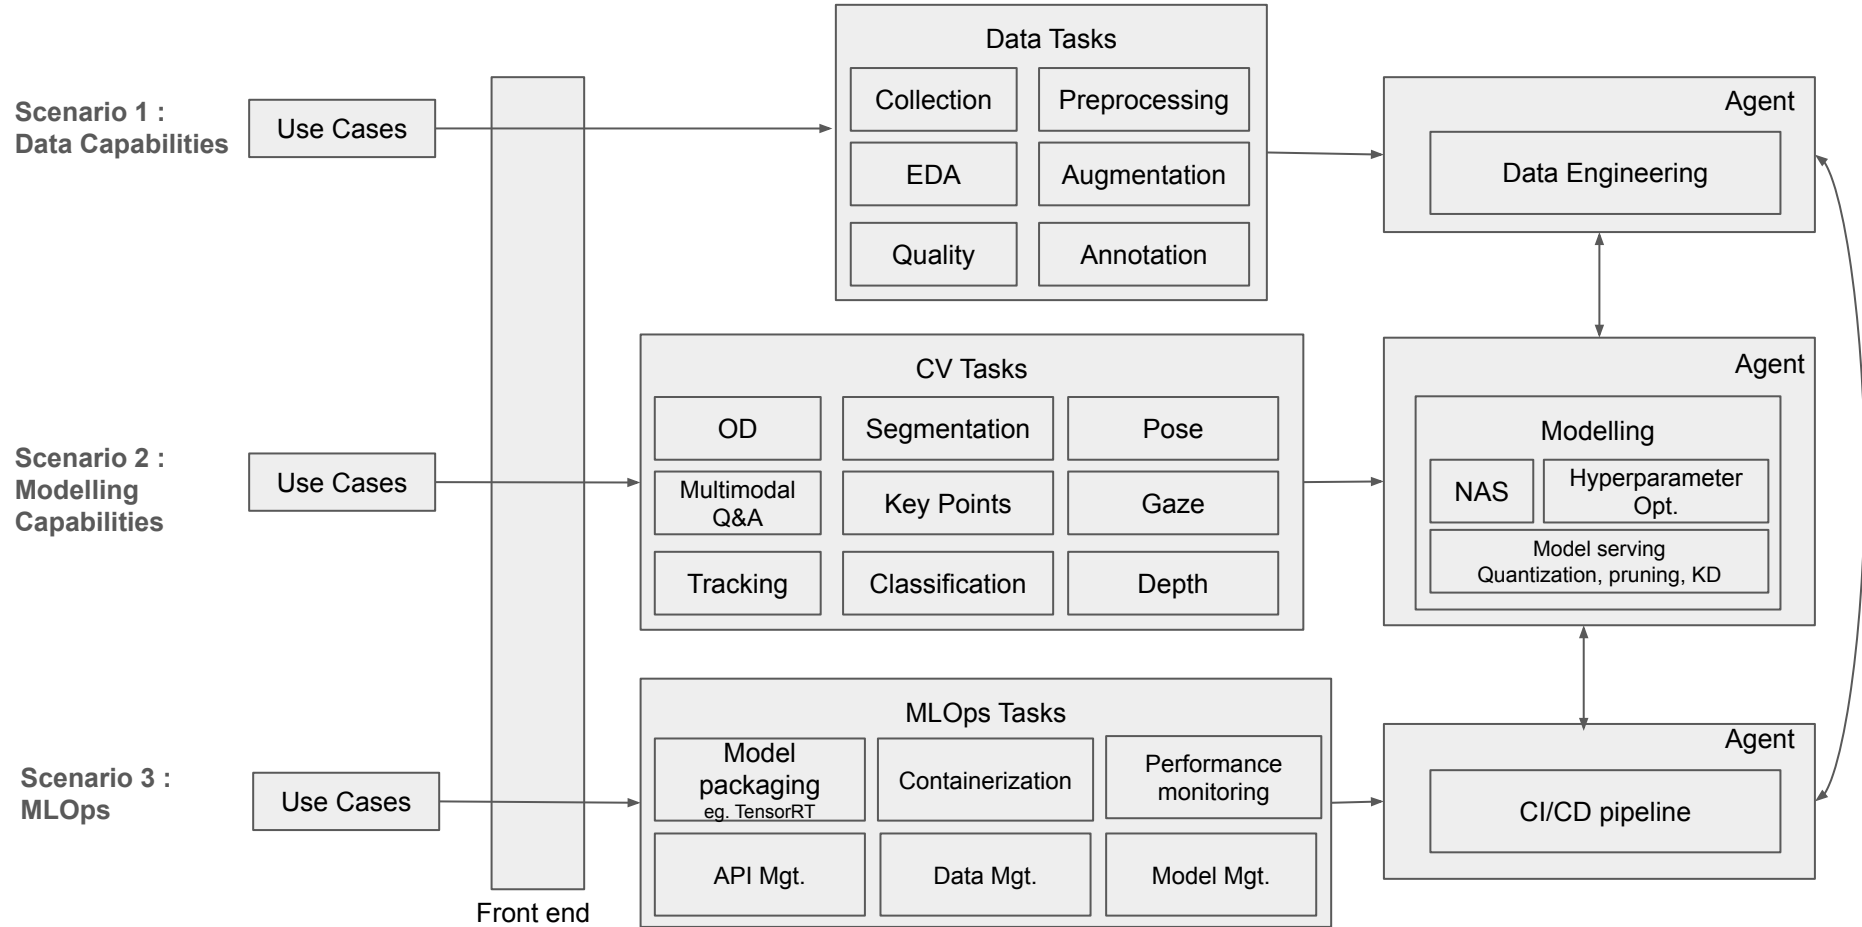
Model challenges

- HPO can get complex if model size and search space is high.
- HPO not applicable for foundation models. Might only need efficient PEFT.
- Cost vs Performance tradeoffs


Deployment challenges


- Don't use proper logging for automatic error correction
- No human in the loop and/or unit testing of production level code
- Efficient serving using quantization, pruning and model conversion (ONNX Format) for different hardware backends
- Security and privacy risks, will require agent to run in a sandboxed environment (eg. docker)

Design Approach : How use cases trickle down to the tech stack

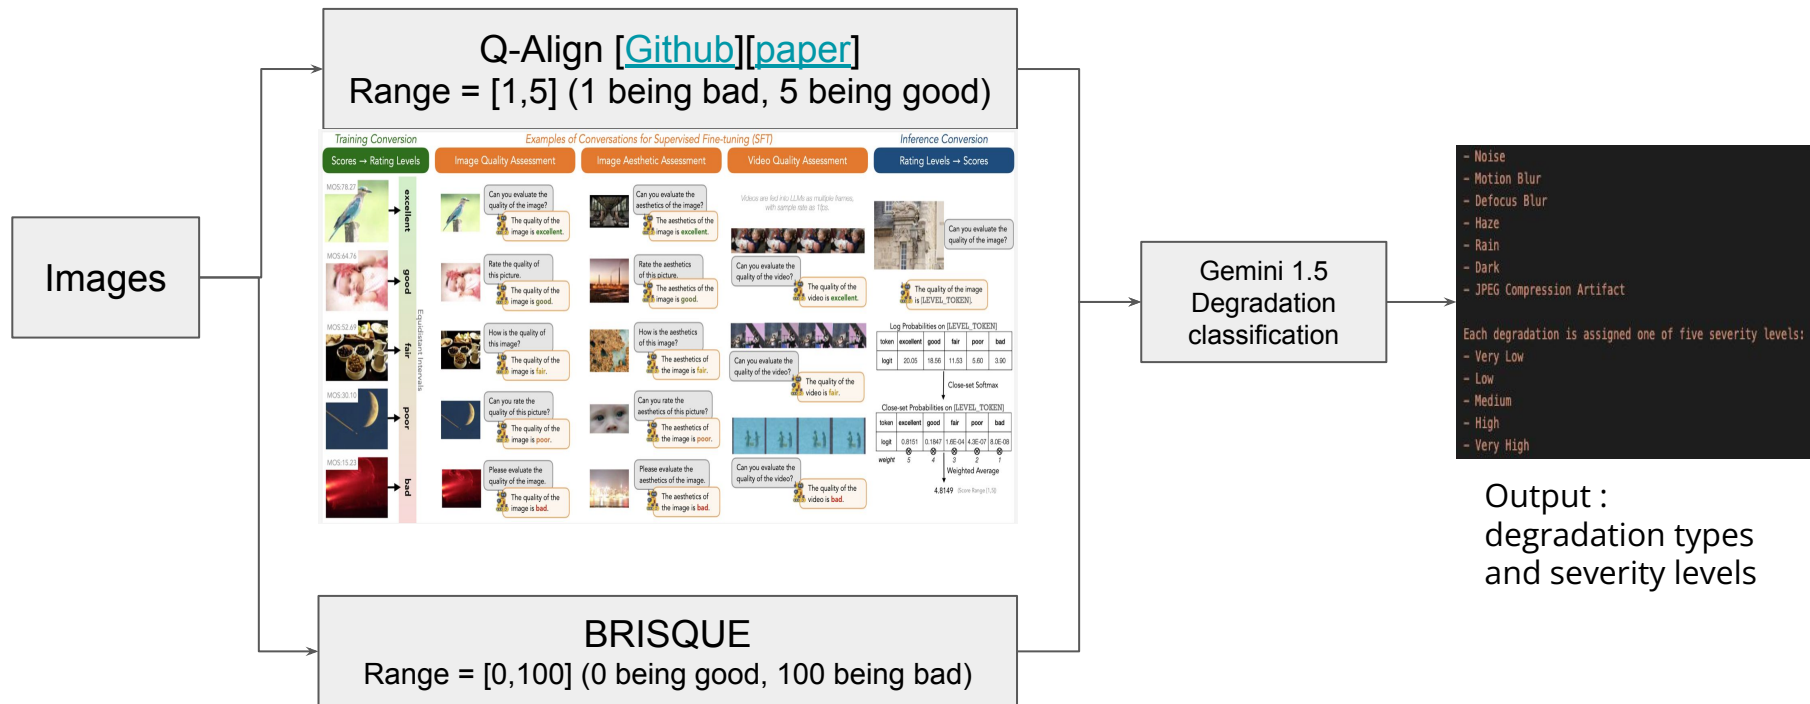


Quality assessment agent : NR-IQA

NR Method	Model names	Description
Q-Align 	qalign (with quality[default], aesthetic options)	Large vision-language models
QualiCLIP(+)	qualiclip, qualiclip+, qualiclip+-clive, qualiclip+-flive, qualiclip+-spaq	QualiCLIP(+) with different datasets, koniq by default
LIQE	liqe, liqe_mix	CLIP based method
ARNIQA	arniqa, arniqa-live, arniqa-csiq, arniqa-tid, arniqa-kadid, arniqa-clive, arniqa-flive, arniqa-spaq	ARNIQA with different datasets, koniq by default
TOPIQ	topiq_nr, topiq_nr-flive, topiq_nr-spaq	TOPIQ with different datasets, koniq by default
TReS	tres, tres-flive	TReS with different datasets, koniq by default
FID	fid	Statistic distance between two datasets
CLIPQA(+)	clipiqa, clipiqa+, clipiqa+_vitL14_512, clipiqa+_rn50_512	CLIPQA(+) with different backbone, RN50 by default
MANIQA	maniqa, maniqa-kadid, maniqa-pipal	MUSIQ with different datasets, koniq by default
MUSIQ	musiq, musiq-spaq, musiq-paq2piq, musiq-ava	MUSIQ with different datasets, koniq by default
DBCNN	dbcnn	

PaQ-2-PiQ	paq2piq	
HyperIQA	hyperiqa	
NIMA	nima, nima-vgg16-ava	Aesthetic metric trained with AVA dataset
WaDIQaM	wadiqam_nr	
CNNIQA	cnniqa	
NRQM(Ma) ²	nrqm	No backward
PI(Perceptual Index)	pi	No backward
BRISQUE 	risque, brisque_matlab	No backward
ILNIQE	lniqe	No backward
NIQE	niqe, niqe_matlab	No backward
PIQE	piqe	No backward

Quality Assessment agent : Degradation classifier



Data preprocessing - auto

